

Spotify Analysis

(Red) Shift Happens

Kelsey Li, Gloria Majchrzak, Stevan Thomas



Agenda:

Introduction / Exploration

Descriptive Analytics

Predictive Analytics



Spotify Dataset

Original:

- Audio features of tracks
- Track count: 586,672 tracks
- Track features (variables): 20
- Time frame: 1922-2021

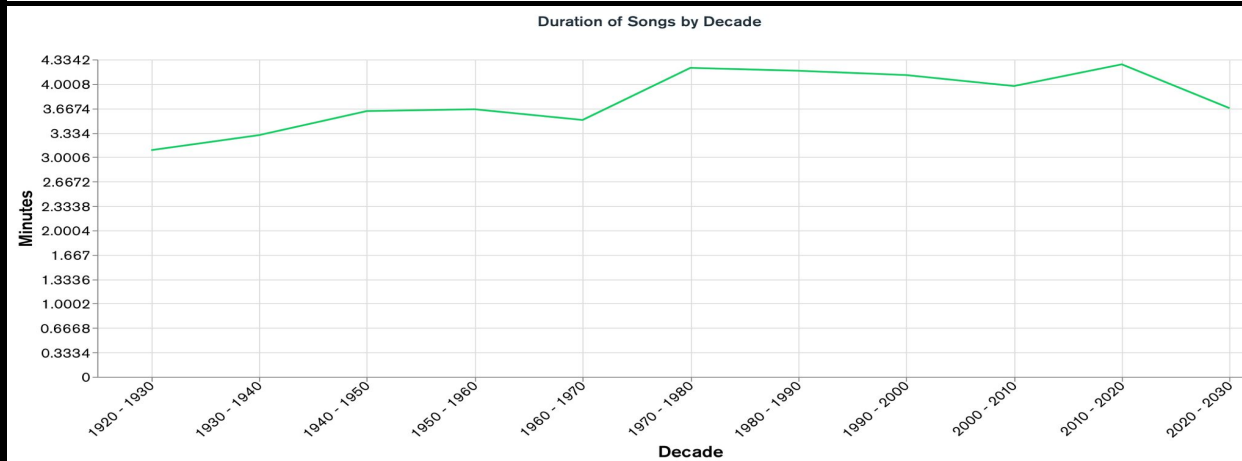
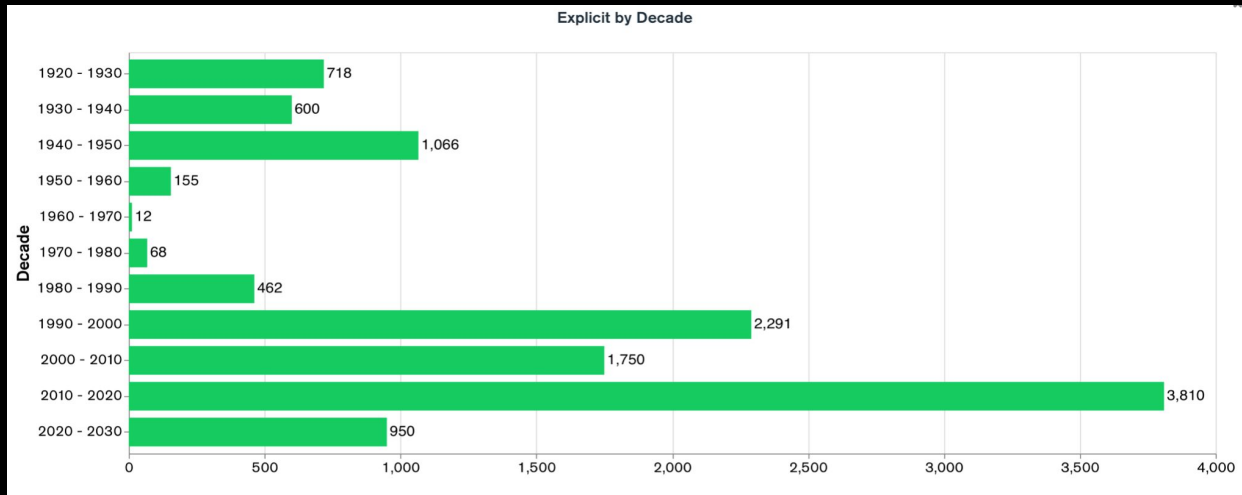
Prepared:

- Decade: created from release_date
- Artist1, Artist2, Artist3, Artist4: created by splitting list in artist variable
- Is_popular: binary created from cutoff of popularity variable
- N_artists: count of artists for each track

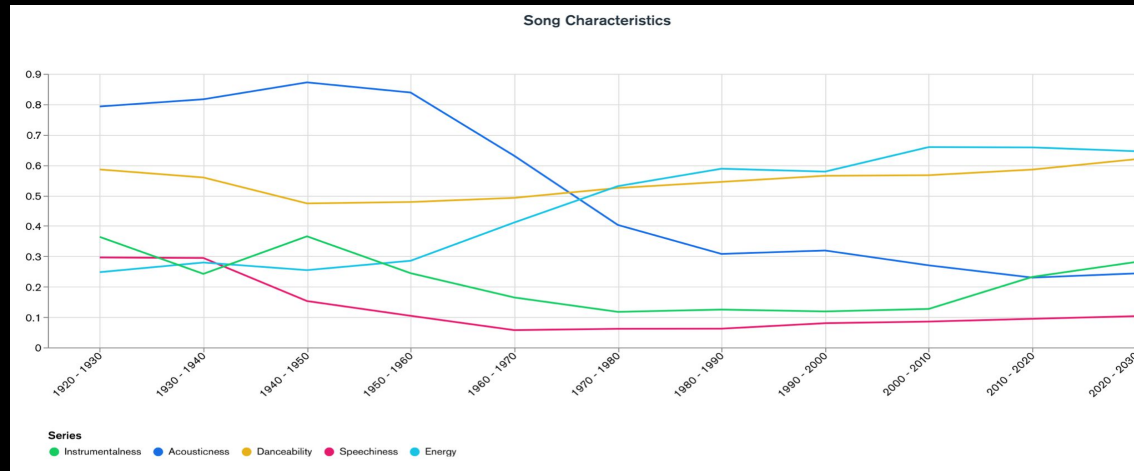
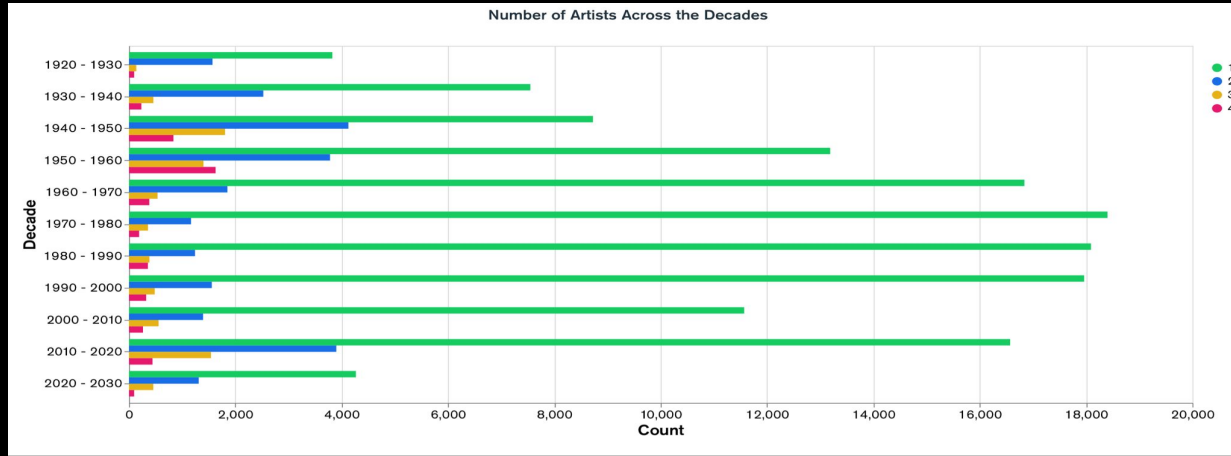
Data Type	Features
Primary	id (generated by Spotify)
Numeric	Acousticness, danceability, energy, duration_ms, instrumentalness, valence, popularity, tempo, liveness, loudness, speechiness
Dummy (0,1)	Mode, explicit
Categorical	Key, time signature, artists, id_artists, release_date, name



Descriptive Analytics



Descriptive Analytics cont.



Predictive Analytics - Multinomial Logistic Regression

Model 1:

- Factorize decade variable
- 15 numeric variables
- Data split: 75-25
- Accuracy: 0.3106

Model 2: Model 1 + Interaction

- Factorize decade variable
- 15 numeric variables
- 6 interaction terms
- Accuracy: 0.3157

```
log_model2 <- multinom_reg() %>%  
  set_engine("glmnet") %>%  
  set_mode("classification") %>%  
  translate()  
  
log_two <- log_model2 %>%  
  fit(decade ~ . + explicit*tempo + danceability*energy + duration_ms*speechiness + loudness*energy +  
  mode*key + acousticness*instrumentalness, data=df_train)
```



Model 3: XGBoost

```
[ ] %%R
xg_model<- xgboost(data = train_spot_matrix, label = as.numeric(train.y$class_numeric), max.depth = 6, eta = 1 , nthread = 4, lambda = 1,
  nrounds = 1000, num_class = 11,
  objective = "multi:softmax", eval_metric="mlogloss", prediction = TRUE )
```

Model 4: Neural Networks

```
[12] # Defining the Neural Network Model
```

```
def baseline_model():
    model = Sequential()
    model.add(Dense(3,input_dim = 16, activation = "relu"))
    model.add(Dense(11,activation = "softmax"))
    model.compile(loss = "categorical_crossentropy", optimizer = "adam", metrics = ['accuracy'])
    return model
```



```
#Evaluating the model with k-fold
```

```
kfold = KFold(n_splits = 10, shuffle = True)
results = cross_val_score(estimator, X, dummy_Y, cv = kfold)
print("Baseline: %.2f%% (%.2f%%)" % (results.mean()*100, results.std()*100))
```

```
Baseline: 18.56% (0.11%)
```

